

# Lab 9: The Discrete Cosine Transform and JPEG

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
import scipy as scp
```

```
In [3]: # Define any utility functions (i.e. inner product, FFT). Feel free to reuse as
def inner_prod_2D(x, y):
    # this does the same thing but vectorized and faster.
    return np.sum(x * np.conj(y))

    # N = np.shape(x)[0]
    # inner_prod = 0
    # for i in range(N):
    #     for j in range(N):
    #         inner_prod += x[i,j] * np.conj(y[i,j])
    # return inner_prod

def complex_exp_2D(N, k, l):
    """
    k, l are frequencies. k is frequency in x direction, l is y.
    N is the duration of the signal
    returns three matrices of size NxN, the first containing the complex values of
    e_kl, NN. the second is the real part; the third is the imaginary part.
    """
    x_tiles = np.tile(np.arange(N), (N, 1))
    y_tiles = np.tile(np.arange(N), (N, 1)).T
    partial_1 = k * x_tiles / N
    partial_2 = l * y_tiles / N
    cexp = 1/N * np.exp(-2j * np.pi * (partial_1 + partial_2))

    cexp_real = np.real(cexp)
    cexp_imag = np.imag(cexp)
    return cexp, cexp_real, cexp_imag

def compute_DFT_2D(x):
    # X is a 2D array containing the coefficients of the 2D DFT of the NxN dimensi
    N = np.shape(x)[0]
    X = np.zeros([N, N], dtype=np.complex128)
    for k in range(N):
        for l in range(N):
            cexp, _, _ = complex_exp_2D(N, k, l)
            X[k, l] = inner_prod_2D(x, cexp)
    return X

def compute_iDFT_2D(X):
    # x is an NxN dimensional signal whose DFT is given by X
    # X has dimensions of... NxN
    # it goes from 0 to N-1 along each dimension.
    N = np.shape(X)[0]
    x = np.zeros((N, N), dtype=np.complex128)

    for m in range(N):
```

```

for n in range(N):
    sum_val = 0

    # Process the first half of frequencies
    for k in range(N//2 + 1): # 0 to N/2
        for l in range(N):
            exponent = 2j * np.pi * (m * k / N + n * l / N)
            current_term = X[k, l] * np.exp(exponent)

            # Add contribution from the conjugate pair
            # Skip for k=0 (which are their own conjugates)
            if k > 0:
                conj_k = (N - k) % N
                conj_l = (N - l) % N
                conj_exponent = 2j * np.pi * (m * conj_k / N + n * conj_l / N)
                current_term += np.conj(X[k, l]) * np.exp(conj_exponent)

            sum_val += current_term

    # convert coordinate system to match original
    x[N-m if m != 0 else m, N-n if n != 0 else n] = (sum_val / N)

return x.T

```

# 1. Image Compression

## 1.1 DCT in Two Dimensions

```

In [4]: def discrete_cosine(k, l, N):
        x_tiles = np.tile(np.arange(N), (N, 1))
        x_tiles = 2*x_tiles + 1
        y_tiles = np.tile(np.arange(N), (N, 1)).T
        y_tiles = 2*y_tiles + 1
        partial_1 = k * x_tiles / (2*N)
        partial_2 = l * y_tiles / (2*N)
        cos = np.cos(np.pi * partial_1)*np.cos(np.pi * partial_2)
        return cos

```

```

In [5]: def compute_2D_DCT(x):
        """
        X = scp.fft.dct(scp.fft.dct(x, axis=0, norm='ortho'), axis=1, norm='ortho')
        return X
        """
        N = np.shape(x)[0]
        X = np.zeros([N, N], dtype=np.complex128)
        for k in range(N):
            for l in range(N):
                cos = discrete_cosine(k, l, N)
                if k == 0:
                    c1 = 1/np.sqrt(2)
                else:
                    c1 = 1
                if l == 0:
                    c2 = 1/np.sqrt(2)
                else:
                    c2 = 1

```

```

        X[k, l] = inner_prod_2D(x, c1*c2*cos)*(2/N)
    return X

```

```

In [6]: def compute_2D_iDCT(X):
    '''
    x = scp.fft.idct(scp.fft.idct(X, axis=0, norm='ortho'), axis=1, norm='ortho')
    return x
    N = np.shape(X)[0]
    Xs = np.zeros([N, N], dtype=np.complex128)
    for k in range(N):
        for l in range(N):
            if k == 0:
                alpha_k = 1/np.sqrt(2)
            else:
                alpha_k = 1
            Xs[k, l] = inner_prod_2D(X, (alpha_k**2)*discrete_cosine(k, l, N).T*discrete_cosine(l, k, N))
    return Xs*(2/N)
    '''

    N = np.shape(X)[0]
    x = np.zeros([N, N], dtype=np.complex128)
    normalization_factor = 2/N
    for m in range(N):
        for n in range(N):
            for k in range(N):
                for l in range(N):
                    if k == 0:
                        c1 = 1/np.sqrt(2)
                    else:
                        c1 = 1
                    if l == 0:
                        c2 = 1/np.sqrt(2)
                    else:
                        c2 = 1

                    x[m,n] += X[k, l] * c1 * c2 * np.cos(np.pi * k * (2*m+1)/(2*N)) * np.cos(np.pi * l * (2*n+1)/(2*N))

    return x.T * normalization_factor

```

```

In [ ]: # If your functions run too slow, consider using scp.fft.dct

```

## 1.2 Image Compression

```

In [7]: def compress_block_DCT(x: np.ndarray, K: int):
    X = compute_2D_DCT(x)
    X_magnitude = np.abs(X)
    threshold = np.sort(X_magnitude.flatten())[-K]
    X[X_magnitude<=threshold] = 0
    return compute_2D_iDCT(X)

def compress_block_DFT(x, K):
    # Implement function to compress an 8x8 block using DFT
    X = compute_DFT_2D(x)
    X_magnitude = np.abs(X)
    threshold = np.sort(X_magnitude.flatten())[-K]

```

```
X[X_magnitude<threshold] = 0
return compute_idFT_2D(X)
```

```
In [8]: def compress_image_DCT(x, K):
        H, W = x.shape # Get dimensions of the input array
        blocks = []

        for i in range(0, H, 8):
            for j in range(0, W, 8):
                blocks.append(x[i:i+8, j:j+8])

        for i in range(len(blocks)):
            blocks[i] = compress_block_DCT(blocks[i], K)
        new_array = np.zeros((H, W)) # Create empty array
        index = 0
        for i in range(0, H, 8):
            for j in range(0, W, 8):
                new_array[i:i+8, j:j+8] = blocks[index]
                index += 1
        return np.abs((new_array.real))

def compress_image_DFT(x, K):
    # Implement function to partition an image into 8x8 blocks and compress the bl
    H, W = x.shape # Get dimensions of the input array
    blocks = []

    for i in range(0, H, 8):
        for j in range(0, W, 8):
            blocks.append(x[i:i+8, j:j+8])



    for i in range(len(blocks)):
        blocks[i] = compress_block_DFT(blocks[i], K)
    new_array = np.zeros((H, W)) # Create empty array
    index = 0
    for i in range(0, H, 8):
        for j in range(0, W, 8):
            new_array[i:i+8, j:j+8] = blocks[index]
            index += 1
    return new_array.real
```



```
In [10]: imgB_prenoise = plt.imread("imgB_prenoise.png")
```

```
In [11]: # Try these functions out on the provided sample image A for K = 4, 8, 16, 32
plt.figure()
output1 = compress_image_DCT(imgB_prenoise, 4)
plt.imshow(output1, cmap='gray')
plt.figure()
output2 = compress_image_DCT(imgB_prenoise, 8)
plt.imshow(output2, cmap='gray')
plt.figure()
output3 = compress_image_DCT(imgB_prenoise, 16)
plt.imshow(output3, cmap='gray')
plt.figure()
output4 = compress_image_DCT(imgB_prenoise, 32)
plt.imshow(output4, cmap='gray')
```

```
<ipython-input-8-cd5e73dc7f87>:15: ComplexWarning: Casting complex values to real  
discards the imaginary part  
new_array[i:i+8, j:j+8] = blocks[index]
```

Out[11]: <matplotlib.image.AxesImage at 0x7ee17c646790>







```
In [12]: #now for the DFT
plt.figure()
output5 = compress_image_DFT(imgB_prenoise, 4)
plt.imshow(output5, cmap='gray')
plt.figure()
output6 = compress_image_DFT(imgB_prenoise, 8)
plt.imshow(output6, cmap='gray')
plt.figure()
output7 = compress_image_DFT(imgB_prenoise, 16)
```



```
plt.imshow(output7, cmap='gray')
plt.figure()
output8 = compress_image_DFT(imgB_prenoise, 32)
plt.imshow(output8, cmap='gray')
```

```
<ipython-input-8-cd5e73dc7f87>:34: ComplexWarning: Casting complex values to real
discards the imaginary part
  new_array[i:i+8, j:j+8] = blocks[index]
```

Out[12]: <matplotlib.image.AxesImage at 0x7ee16478b710>







```
[18, 22, 37, 56, 68, 109, 103, 77],
[24, 36, 55, 64, 81, 104, 113, 92],
[49, 64, 78, 87, 103, 121, 120, 101],
[72, 92, 95, 98, 112, 100, 103, 99]])
```

```
In [14]: def quantize(Q, block):
# Implement function to quantize an 8x8 block
#print(block)
return np.round(block/Q)

def dequantize(Q, block):
# Implement function to dequantize an 8x8 block
return block*Q
```

```
In [15]: def quantize_image(Q, x):
# Implement function to quantize an image
H, W = x.shape # Get dimensions of the input array
blocks = []
for i in range(0, H, 8):
    for j in range(0, W, 8):
        blocks.append(x[i:i+8, j:j+8])

for i in range(len(blocks)):
    blocks[i] = quantize(Q,blocks[i])
new_array = np.zeros((H, W)) # Create empty array
index = 0
for i in range(0, H, 8):
    for j in range(0, W, 8):
        new_array[i:i+8, j:j+8] = blocks[index]
        index += 1
return np.abs((new_array.real))


def dequantize_image(Q, x):
# Implement function to dequantize an image
H, W = x.shape # Get dimensions of the input array
blocks = []

for i in range(0, H, 8):
    for j in range(0, W, 8):
        blocks.append(x[i:i+8, j:j+8])

for i in range(len(blocks)):
    blocks[i] = dequantize(Q,blocks[i])
new_array = np.zeros((H, W)) # Create empty array
index = 0
for i in range(0, H, 8):
    for j in range(0, W, 8):
        new_array[i:i+8, j:j+8] = blocks[index]
        index += 1
return np.abs((new_array.real))
```

```
In [16]: imgB_prenoise = plt.imread("imgB_prenoise.png")
plt.figure()
output10 = quantize_image(Q,imgB_prenoise*255)
output10 = dequantize_image(Q,output10)
output10 = output10/255
plt.imshow(output10, cmap='gray')
```

```
Out[16]: <matplotlib.image.AxesImage at 0x7ee16460a550>
```



## 1.4 Image Reconstruction

```
In [17]: # Load image
img = plt.imread('imgB_prenoise.png')
```

```
In [18]: def get_reconstruction_error(x_bar, x):
# Calculate distance between x_bar and x
return np.linalg.norm(x_bar - x)
```

```
In [19]: # Compress and reconstruct using functions from 1.2
# Make plots of reconstructions for K = 4, 8, 16, 32
# Calculate reconstruction error for different values of K, and plot
#DFT
K_values = [4, 8, 16, 32]
DFT_outputs = [output5, output6, output7, output8]
errors = []
for k, x_bar in zip(K_values, DFT_outputs):
    error = get_reconstruction_error(x_bar, img)
    errors.append(error)
plt.figure()
plt.plot(K_values, errors, marker='o', linestyle='--', color='b', label="Reconstr


# Labels and Title
plt.xlabel("K Values")
plt.ylabel("Reconstruction Error")
plt.title("Reconstruction Error vs. K")
plt.grid(True, linestyle="--", alpha=0.6)
plt.legend()

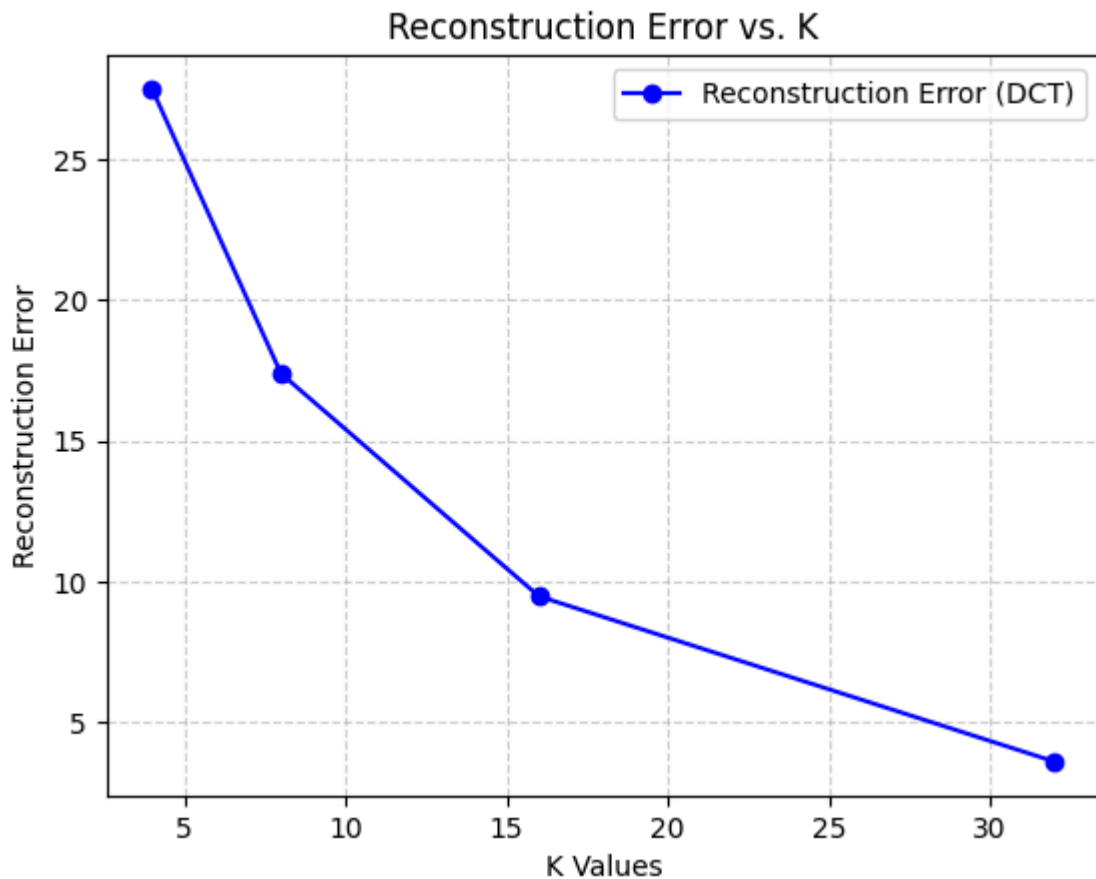
# Show the plot
plt.show()
```

```
#####
#DCT
K_values = [4, 8, 16, 32]
DCT_outputs = [output1, output2, output3, output4]
errors = []
for k, x_bar in zip(K_values, DCT_outputs):
    error = get_reconstruction_error(x_bar, img)
    errors.append(error)
plt.figure()
plt.plot(K_values, errors, marker='o', linestyle='-', color='b', label="Reconstr

# Labels and Title
plt.xlabel("K Values")
plt.ylabel("Reconstruction Error")
plt.title("Reconstruction Error vs. K")
plt.grid(True, linestyle="--", alpha=0.6)
plt.legend()

# Show the plot
plt.show()
```






**Comments:** As we can see the more we compress signals the more we lose information which is a general trend in all compression schemes.

```
In [20]: # Compress and reconstruct using functions from 1.3 (quantization)
# Make plots of reconstructions using different values of the quantization matrix
Q1 = np.array([[9, 3, 4, 5, 12, 11, 8, 6],
               [11, 15, 4, 3, 1, 9, 8, 7],
               [14, 13, 6, 4, 4, 7, 9, 6],
               [4, 7, 2, 9, 5, 8, 8, 6],
               [8, 2, 7, 6, 8, 9, 13, 7],
               [4, 6, 5, 6, 8, 4, 13, 9],
               [9, 6, 8, 8, 13, 21, 12, 11],
               [7, 9, 9, 8, 12, 10, 13, 9]])

plt.figure()
output11 = quantize_image(Q1,img*255)
output11 = dequantize_image(Q1,output11)
output11 = output11/255
plt.imshow(output11, cmap='gray')
```

Out[20]: <matplotlib.image.AxesImage at 0x7ee17c547a10>

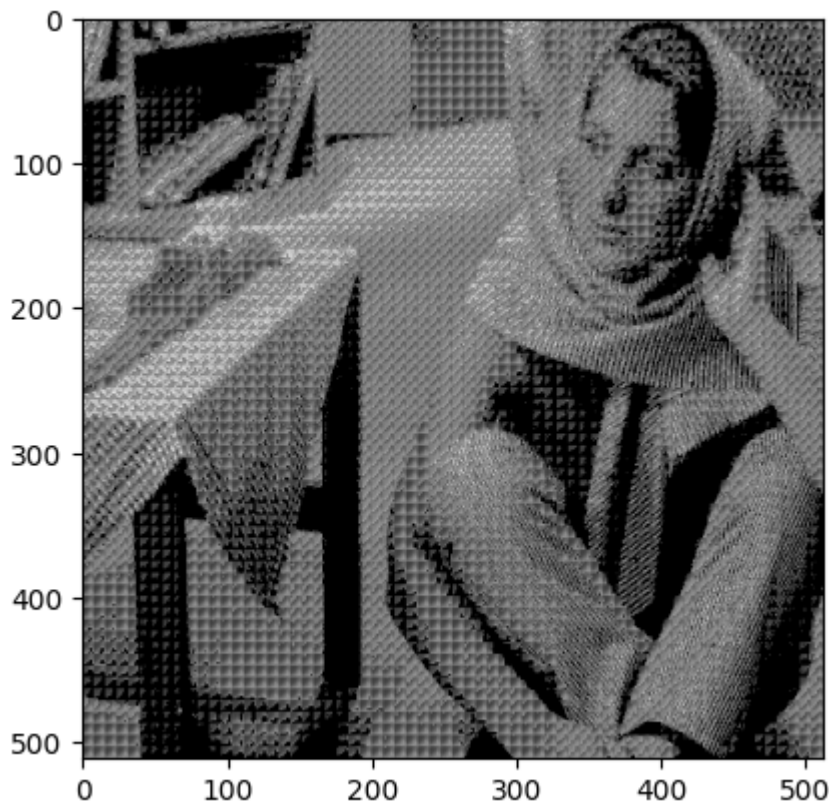


Comment: Here as we can see we have very low Q that does a decent compression (maybe lowers the size by a factor of two) while also providing a very detailed picture with nearly no lose of information.

```
In [21]: Q2 = np.array([[16, 11, 10, 16, 24, 40, 51, 61],
                        [12, 12, 14, 19, 26, 58, 60, 55],
                        [14, 13, 16, 24, 40, 57, 69, 56],
                        [14, 17, 22, 29, 51, 87, 80, 62],
                        [18, 22, 37, 56, 68, 109, 103, 77],
                        [24, 36, 55, 64, 81, 104, 113, 92],
                        [49, 64, 78, 87, 103, 121, 120, 101],
                        [72, 92, 95, 98, 112, 100, 103, 99]])

Q2 = Q2+70
plt.figure()
output12 = quantize_image(Q2,img*255)
output12 = dequantize_image(Q2,output12)
output12 = output12/255
plt.imshow(output12, cmap='gray')
```

Out[21]: <matplotlib.image.AxesImage at 0x7ee17c54e490>




Comment: For moderately high Q (that doesn't destroy information) you get decent amount of compression while also preserving the information, but as we can see the black dots got blacker and we lost a lot of shades in the process of compression.

```
In [22]: Q3 = np.array([[16, 11, 10, 16, 24, 40, 51, 61],
                        [12, 12, 14, 19, 26, 58, 60, 55],
                        [14, 13, 16, 24, 40, 57, 69, 56],
                        [14, 17, 22, 29, 51, 87, 80, 62],
                        [18, 22, 37, 56, 68, 109, 103, 77],
                        [24, 36, 55, 64, 81, 104, 113, 92],
                        [49, 64, 78, 87, 103, 121, 120, 101],
                        [72, 92, 95, 98, 112, 100, 103, 99]])

Q3 = Q3+700
plt.figure()
output13 = quantize_image(Q3,img*255)
output13 = dequantize_image(Q3,output13)
output13 = output13/255
plt.imshow(output13, cmap='gray')
```

```
Out[22]: <matplotlib.image.AxesImage at 0x7ee164593910>
```





*\*\* italicized text*Comments:\*\* As we can see high values of Q will result in more compression but also you will lose much more information and if the Q is high enough you will lose all information.

A rule that we can deduce is High Q-> more compression -> less information and vice versa.